REPORT ON TRANSFORMER ARCHITECTURE IN NATURAL LANGUAGE PROCESSING

INDEX

INTRODUCTION

Natural Language Processing (NLP) is a branch of Artificial Intelligence (AI) that enables computers to understand, interpret, and generate human language. NLP is widely used in applications such as machine translation, chatbots, voice assistants, sentiment analysis, text summarization, and question-answering systems.

Earlier NLP models mainly used Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks. Although these models were effective, they processed words one after another in sequence. This made training slow and created difficulties in remembering information from long sentences.

To overcome these limitations, researchers introduced the **Transformer architecture** in 2017 through the research paper **"Attention Is All You Need."** Instead of processing words sequentially, Transformers process all words at the same time using a mechanism called **Attention**. This allows the model to understand the relationship between words regardless of how far apart they are in a sentence.

The Transformer architecture has become the foundation of many modern AI models such as BERT, GPT, T5, RoBERTa, and many others. These models achieve high accuracy in language understanding and generation tasks.

The main components of a Transformer are Word Embeddings, Positional Encoding, Self-Attention, Multi-Head Attention, Encoder, Decoder, Feed Forward Networks, and Softmax. Each component plays an important role in helping the model understand and process natural language efficiently.

WORD EMBEDDINGS AND POSITIONAL ENCODING

Before a computer can understand language, words must first be converted into numerical values. This process is known as **Word Embedding**.

A word embedding represents every word as a vector of numbers. Words with similar meanings usually have similar vector representations. This allows the model to capture semantic relationships between words.

For example, consider the sentence:

"Artificial Intelligence is powerful."

After tokenization, the sentence becomes:

- Artificial
- Intelligence
- is
- powerful

Each word is converted into a numerical vector such as:

- Artificial → [0.45, 0.82, 0.16]
- Intelligence → [0.71, 0.39, 0.55]
- is → [0.12, 0.08, 0.44]
- powerful → [0.87, 0.91, 0.22]

These vectors are called **Word Embeddings**.

Why is Positional Encoding Needed?

Unlike RNNs, Transformers process all words simultaneously. Therefore, they cannot naturally determine the order of words.

For example:

- "The cat chased the dog."
- "The dog chased the cat."

Both sentences contain the same words but have different meanings because of the order of the words.

To solve this problem, **Positional Encoding** is added to each word embedding. Positional Encoding provides information about the position of each word in the sentence, allowing the Transformer to understand the sequence correctly.

Thus, Word Embeddings provide the meaning of words, while Positional Encoding provides their positions.

SELF-ATTENTION MECHANISM

The most important innovation of the Transformer architecture is the **Self-Attention Mechanism**.

Self-Attention allows each word to focus on other important words in the same sentence while understanding its meaning.

Consider the sentence:

"The animal didn't cross the road because it was tired."

The word **"it"** refers to **"animal."**

A traditional model may struggle to identify this relationship because the words are separated by several other words. However, the Self-Attention mechanism correctly identifies that "it" refers to "animal."

Query, Key, and Value

Each word is transformed into three vectors:

- Query (Q)
- Key (K)
- Value (V)

These vectors help determine how much attention one word should pay to another.

The attention score is calculated by comparing the Query with all Keys. The scores indicate the importance of each word. Higher scores mean greater importance.

After calculating the attention scores, the model combines the corresponding Value vectors to produce a new representation for each word.

This process enables the Transformer to understand context much more effectively than earlier models.

Advantages of Self-Attention

- Captures long-distance relationships between words.
- Processes all words simultaneously.
- Improves contextual understanding.
- Increases model accuracy.
- Reduces training time.

Because of these advantages, Self-Attention has become one of the most important concepts in modern Natural Language Processing.

MULTI-HEAD ATTENTION AND SOFTMAX

Multi-Head Attention

Instead of calculating attention only once, the Transformer performs multiple attention operations simultaneously. Each attention operation is called an **Attention Head**.

Different heads focus on different aspects of the sentence.

For example:

- One head focuses on grammar.
- Another head focuses on word meaning.
- Another head focuses on sentence context.
- Another head identifies relationships between words.

After all attention heads complete their calculations, their outputs are combined to produce a richer and more informative representation.

Because multiple heads work together, the model learns more information than a single attention mechanism.

Softmax Function

After calculating the attention scores, the Transformer uses the **Softmax Function**.

The Softmax function converts numerical scores into probabilities.

For example:

Attention Scores:

- 5.2
- 3.4
- 1.1

After applying Softmax:

- 0.80
- 0.15

- 0.05

The probabilities always add up to 1.

The word with the highest probability receives the greatest attention during processing.

The Softmax function helps the model decide which words are most important while generating the output.

Together, Multi-Head Attention and Softmax enable the Transformer to understand language accurately and efficiently.

ENCODER AND DECODER ARCHITECTURE

The Transformer architecture consists of two major parts:

- Encoder
- Decoder

Encoder

The Encoder receives the input sentence and converts it into meaningful numerical representations.

Each Encoder layer consists of:

- Multi-Head Attention
- Add and Normalize
- Feed Forward Neural Network
- Add and Normalize

The Transformer usually contains several Encoder layers stacked one above another. Each layer improves the understanding of the sentence.

Decoder

The Decoder generates the output sentence one word at a time.

Each Decoder layer contains:

- Masked Multi-Head Attention
- Encoder–Decoder Attention
- Feed Forward Network
- Linear Layer
- Softmax Layer

The Decoder uses the information produced by the Encoder together with the previously generated words to predict the next word.

For example, during language translation, the Encoder processes the English sentence while the Decoder generates the translated sentence in another language.

The combination of Encoder and Decoder enables the Transformer to perform tasks such as translation, summarization, and text generation with high accuracy.

APPLICATIONS OF TRANSFORMERS

The Transformer architecture has become the foundation of many Artificial Intelligence applications.

Some important applications include:

Machine Translation

Transformers translate text from one language to another with high accuracy.

Example:

English → French

Chatbots

Modern chatbots use Transformers to understand user questions and generate meaningful responses.

Text Summarization

Transformers can generate short summaries from long articles while preserving important information.

Sentiment Analysis

Companies use Transformers to determine whether customer reviews are positive, negative, or neutral.

Question Answering

Transformers answer questions by understanding the context of documents or conversations.

Speech Recognition

Voice assistants convert spoken language into text using Transformer-based models.

Text Generation

Language models such as GPT generate essays, stories, emails, and programming code using the Transformer architecture.

Because of their high accuracy and efficiency, Transformers are widely used in education, healthcare, finance, business, research, and many other fields.

CONCLUSION

The Transformer architecture has brought a revolutionary change in the field of Natural Language Processing. Unlike traditional sequence-based models, Transformers process all words simultaneously using the Self-Attention mechanism. This makes them faster, more efficient, and capable of understanding long-range relationships between words.

Important components such as Word Embeddings, Positional Encoding, Self-Attention, Multi-Head Attention, Encoder, Decoder, and Softmax work together to understand and generate human language effectively.

Today, almost every advanced language model, including GPT, BERT, T5, and many others, is built upon the Transformer architecture. As Artificial Intelligence continues to evolve, Transformers are expected to play an even greater role in developing intelligent systems for communication, healthcare, education, business, and scientific research.

The Transformer architecture is therefore considered one of the most significant breakthroughs in modern Artificial Intelligence and Natural Language Processing.